

Overload as a Cusp Catastrophe: Identifiability Limits and a Ground-Truth Failure of Early-Warning Indicators in Wearable Physiology

Neil Thomas Mathew 

Department of Computer Applications

CHRIST (Deemed to be University), Bengaluru, India

ORCID: 0009-0001-8802-7376

Abstract—Sensory and executive overload is usually modelled by labelling states and counting transitions between them. Models of that kind fit any data and predict nothing that could refute them. We derive overload instead from the geometry of a cusp catastrophe. A latent load coordinate relaxes in a quartic potential whose asymmetry follows momentary demand and whose bistability follows a slow recovery-debt variable, so the discrete functional states become metastable basins and the transition kernel follows from Kramers escape rates. Two exponents are fixed in advance: hysteresis width grows as the $3/2$ power of the splitting factor, and near the fold the variance and log-autocorrelation scale as $\mp 1/2$ in distance to the tipping point. Conditional on the slow variable the drift is linear in a reparameterised coefficient vector, so estimation reduces to profiled least squares and is exact. Three results follow, and two of them concern method rather than physiology. First, a Monte Carlo study over 3.6×10^3 simulated series shows that at the series lengths wearable corpora actually supply only the relaxation rate and the diffusion coefficient are identified; the control parameters that carry the geometry are not, which bounds what any cusp analysis of such data can claim. Second, the rolling-window autocorrelation and variance indicators used throughout the early-warning literature return exponents of the wrong sign on data simulated from a model that genuinely contains a fold, so we withdraw our own early-warning test rather than report it. Third, across 147 recordings from 40 people in three open corpora, the one identifiable prediction is not supported above a surrogate-calibrated null. Measured power bounds this to excluding strong and moderate dynamics. The work is motivated by autistic adults; none are analysed here, and a pre-registered validation is specified.

Index Terms—Critical transitions, cusp catastrophe, hysteresis, early-warning signals, identifiability, electrodermal activity, wearable sensing, sensory overload.

I. Introduction

Autistic adults describe functional breakdown under combined sensory and cognitive demand in consistent terms: it arrives suddenly, and it clears slowly [1]. Adult cohort studies couple sensory atypicality to executive difficulty with large effect sizes [2], [3], so the phenomenon is not in doubt. What is missing is a formal account of the dynamics.

Existing models do not supply one. Computational work on autism is dominated by Bayesian and predictive-coding framings whose empirical support a ten-year review found mixed and methodologically inconsistent [4], and the most

recent dynamical treatment of symptom trajectories is explicitly theoretical and collects no data [5]. Applied state-transition work typically assigns states from experimental condition labels and then counts a transition matrix, which is close to circular: the recovered matrix largely re-describes the block structure of the protocol, and yields no prediction that observation could contradict.

We take a different route. Rather than positing states and fitting transitions between them, we posit a one-dimensional stochastic relaxation in a cusp potential and let the phenomenology follow. The load coordinate x_t sits in $V(x; a, b) = \frac{1}{4}x^4 - \frac{1}{2}ax^2 - bx$, where the normal factor b_t carries momentary demand and the splitting factor a_t carries depleted regulatory capacity. The substantive claim is that capacity does not add to load. It changes the geometry. Below a critical a the load tracks demand and returns; above it a second attractor appears, and with it a fold, a hysteresis loop and critical slowing down. Sudden onset and slow recovery stop being observations to report and become consequences to check, and because the geometry is closed-form those consequences are numerical exponents that data can reject. This paper is about how far that promise survives contact with real wearable data. Contributions.

- 1) A cusp formulation in which the five functional states are geometric regions rather than labels, and the 5×5 transition kernel is generated by nine interpretable parameters instead of twenty counted frequencies (Sec. II).
- 2) An exact profiled-least-squares estimator, and a high-replication identifiability study showing which parameters are recoverable at the series lengths wearable corpora supply (Sec. IV), which bounds every claim such an analysis can make, including ours.
- 3) A ground-truth demonstration that the standard early-warning indicators report the wrong sign on data containing a genuine fold, plus a measurement of the eightfold size inflation of the nominal test on near-unit-root series (Sec. V).
- 4) A calibrated null on 147 recordings, with power characterised so the result is a bound rather than

an absence.

The target population is autistic adults. No public wearable dataset of that group exists at the temporal density this model needs [6], [7], so none are analysed here, and we say so rather than title the work after a group it does not measure.

II. The Cusp-Hysteretic Model

A. Fast dynamics and a slow debt variable

The latent load obeys a gradient relaxation with additive noise,

$$dx_t = \lambda[-x_t^3 + a_t x_t + b_t]dt + \sigma dW_t, \quad (1)$$

with control parameters linked to observed drivers and to a slow variable A_t that accumulates recovery debt:

$$b_t = \beta_0 + \beta_S S_t + \beta_T T_t + \beta_U U_t, \quad a_t = \alpha_0 + \alpha_A A_t, \quad (2)$$

$$dA_t = \varepsilon[\text{sp}(x_t - x_{\text{ref}}) - A_t]dt, \quad \varepsilon \ll 1. \quad (3)$$

Here sp is a softplus, S_t , T_t and U_t are sensory, task-switch and uncertainty drivers, and λ sets the timescale without touching the geometry. Load above baseline accrues debt, debt raises a_t , and a higher a_t widens the hysteresis loop, so one episode makes the next easier. That is a mechanistic reading of the refractory period qualitative accounts describe.

B. Folds and two scaling laws

Equilibria solve $x^3 - ax - b = 0$. Its discriminant $4a^3 - 27b^2$ is positive exactly when two stable basins coexist. Folds occur where $V''(x) = 0$, that is at $x_{\pm} = \pm\sqrt{a/3}$, giving critical drives $\pm b_c(a)$ with $b_c(a) = \frac{2}{3\sqrt{3}}a^{3/2}$.

Proposition 1 (Hysteresis width, P1). For $a > 0$ the system jumps to the upper branch at $b = +b_c$ and returns at $b = -b_c$, so the loop has width

$$\Delta b(a) = \frac{4}{3\sqrt{3}}a^{3/2}. \quad (4)$$

Linearising (1) about a stable equilibrium gives a relaxation rate $\lambda_{\text{rel}} = 3x^2 - a$, which vanishes at a fold. Critical slowing down is therefore a theorem here rather than an assumption. Writing $\mu = b_c - b$ for the distance to the fold, the normal form gives $\lambda_{\text{rel}} \propto \mu^{1/2}$, hence

Proposition 2 (Early-warning exponents, P2). $\text{Var}(x) \propto \mu^{-1/2}$ and $-\log \text{AC1} \propto \mu^{+1/2}$.

These are point hypotheses about exponents. They differ in kind from the usual directional claim that an indicator rises before a transition, which the sceptical literature shows arises spuriously on almost any autocorrelated series [8], [9].

C. States and transition kernel

Crossing basin membership with the bistability condition partitions the (a, b) plane into five regions (Fig. 1b): calm and focused in the lower basin when monostable, stuck in the lower basin while an overloaded attractor coexists, overloaded in the upper basin, and recovering in the upper basin once demand has fallen below the level that would have put the system there. The last two are the useful ones. Qualitative work uses both words freely, and to our knowledge neither has previously had a precise mathematical meaning. Between-basin moves are noise-induced escapes with Kramers rate $r = \frac{1}{2\pi} \sqrt{|V''(x^*)V''(x_s)|} \exp(-2\Delta V/\sigma^2)$ [10], so the kernel is a deterministic function of (a_t, b_t, σ) : twenty counted probabilities replaced by nine interpretable parameters that move with demand.

III. Materials and Methods

A. Corpora

Three open corpora, each taken from its primary repository rather than a mirror, span a gradient of experimental control: WESAD, 15 subjects in a controlled laboratory protocol [11]; Wearable Exam Stress, 10 students across 30 examinations [12]; and Nurse Stress, 15 nurses across roughly 1250 h of hospital shifts [13]. All supply Empatica E4 electrodermal activity, blood volume pulse, skin temperature and accelerometry. After windowing and quality screening this gives 147 recordings from 40 people, median 194 windows each.

B. Sensor assignment and a negative control

Reading the load coordinate and its drivers from the same sensor would make the model unfalsifiable, since shared measurement noise alone could produce apparent coupling, so the two are assigned to disjoint sensors. The load coordinate of a slow relaxation process must itself be slow, which is measurable rather than a matter of preference: at a 30 s window tonic electrodermal activity has lag-1 autocorrelation 0.98, whereas a cardiac index built from heart rate and RMSSD has 0.19, because RMSSD over so few beats is dominated by beat-detection error. Configuration B (tonic EDA) is therefore primary, configuration C (skin temperature) is a robustness check, and configuration A (the cardiac index) is retained as a negative control in which a working pipeline should find nothing.

C. Estimation

Taking the Euler map of (1) as the model, rather than as an approximation to it, makes the transition density exactly Gaussian and removes discretisation bias. Expanding the drift, the mean increment is linear in a reparameterised coefficient vector,

$$\Delta x_t = \theta_1(-x_t^3) + \theta_2 x_t + \theta_3 A_t x_t + \theta_4 + \theta_5 S_t + \theta_6 T_t + \theta_7 U_t + \eta_t, \quad (5)$$

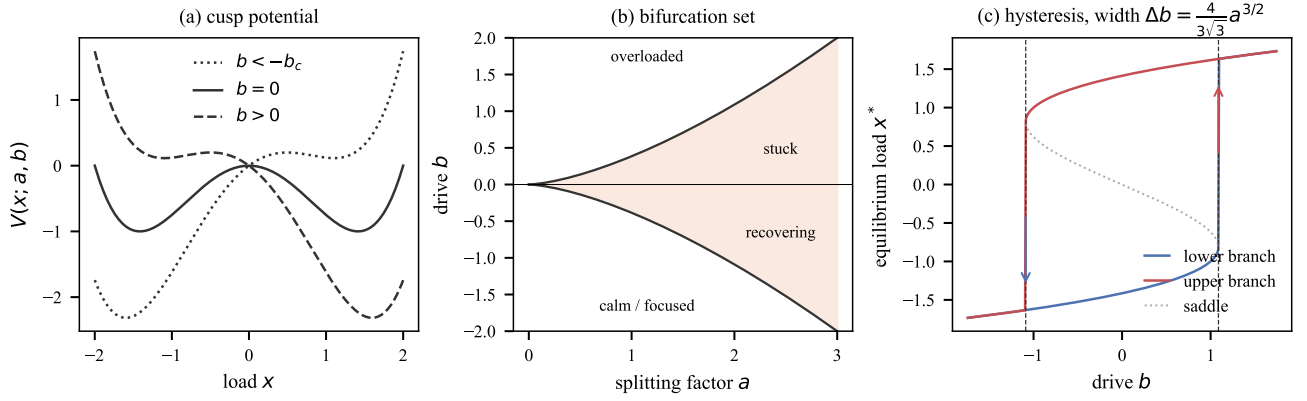


Fig. 1. The geometry from which everything else follows. (a) The cusp potential at three levels of drive b . (b) The bifurcation set in the (a, b) plane, with the five functional states appearing as regions rather than as assigned labels. (c) The hysteresis loop, whose width is fixed by (4) rather than merely being positive.

with $\theta_1 = \lambda$, $\theta_2 = \lambda\alpha_0$, $\theta_3 = \lambda\alpha_A$ and $\theta_{4:7} = \lambda\beta_{0:U}$. Since A_t depends only on ε , the fit is ordinary least squares profiled over a single scalar: exact, with no bounds, no starting values and no local optima. We impose $\theta_1 \geq 0$, because a negative value places the system outside the model rather than merely fitting it badly, and we retain units where that constraint binds so that no reported proportion is conditioned on its own outcome.

D. Nulls

The nominal t -test on θ_1 is not usable here: tonic EDA is close to a unit-root process, and regressing on an integrated series inflates significance in the classical way. Every test is therefore calibrated against surrogate ensembles, both random-walk surrogates matched on length, increment variance and marginal scale, and amplitude-adjusted Fourier surrogates (IAAFT [14]) that preserve spectrum and marginal so only nonlinearity can drive a rejection. Competing models (homogeneous Markov, covariate logistic, Gaussian HMM, the nested Ornstein–Uhlenbeck case, gradient boosting) are scored by held-out one-step predictive log-density.

IV. What This Data Can and Cannot Identify

Before testing anything we asked which parameters are recoverable at the series lengths on offer. This is not a formality. It turned out to be the result that organises the rest of the paper.

We simulated from the model with randomised true parameters and refitted, 600 replicates at each of six lengths, then correlated true against estimated values (Table I, Fig. 2a). At $n = 1000$ most parameters are at least moderately recovered. At $n = 200$, which is the median length in these corpora, only two are: the relaxation rate λ at $r = 0.91$ and the diffusion coefficient σ at $r = 0.95$. The splitting factor α_0 sits at $r = 0.06$, the debt gain α_A at $r = -0.07$, and the slow rate ε at $r = 0.07$.

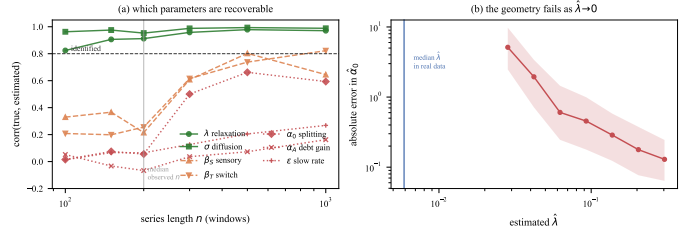


Fig. 2. (a) Recovery correlation against series length, 600 replicates per point. Only λ and σ clear a useful threshold at the median observed length. (b) The failure mechanism: absolute error in the recovered splitting factor rises as $\hat{\lambda}$ falls, because $a = \theta_2/\theta_1$. The median $\hat{\lambda}$ in the real corpora lies below the whole simulated range.

TABLE I

Correlation between true and estimated parameters, 600 replicates per cell. At the median observed length ($n = 200$) only λ and σ are identified. The control parameters carrying the geometry require $n \gtrsim 500$.

Parameter	$n = 100$	$n = 200$	$n = 500$	$n = 1000$
λ (relaxation rate)	0.82	0.91	0.98	0.97
σ (diffusion)	0.96	0.95	0.99	0.99
β_S (sensory drive)	0.33	0.22	0.80	0.65
β_T (task switch)	0.21	0.26	0.74	0.82
α_0 (splitting)	0.02	0.06	0.66	0.59
α_A (debt gain)	0.05	-0.07	0.07	0.16
ε (slow rate)	0.01	0.07	0.21	0.27

The consequence is sharp, and it applies to any study of this kind rather than only to ours. The recovered geometry is built from the ratios $a = \theta_2/\theta_1$ and $b = \theta_4/\theta_1$, so when $\hat{\lambda}$ is small both inflate without bound (Fig. 2b). That is exactly what the real fits show: among units where the constraint does not bind, the median $\hat{\lambda}$ is 0.0059 and the median recovered $\hat{\alpha}_0$ is 12, which would place the attractors near ± 2 on a coordinate standardised to unit variance. Those are not weakly bistable systems. They are systems with no restoring cubic term, whose apparent geometry comes from dividing by a number close to zero.

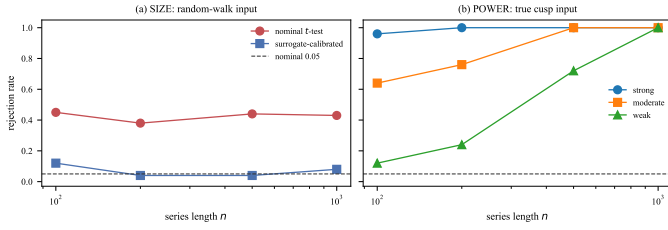


Fig. 3. (a) Size on random-walk input. The nominal test rejects at roughly 0.42 irrespective of series length, while surrogate calibration restores approximately nominal size. (b) Power of the calibrated test on input that genuinely contains a cusp. At the median observed length, strong and moderate dynamics are detected and weak dynamics largely are not.

Note that this median $\hat{\lambda}$ falls below the entire range our simulation covers, so the real recordings sit further into the non-identified regime than any of the 3.6×10^3 synthetic series we used to characterise it.

So exactly one of the five pre-stated predictions is evaluable here: the existence of the restoring cubic term, governed by the one well-identified parameter. P1 needs the ratio of two poorly identified quantities, P2 the fold distance, and P3 to P5 depend on ε and α_A . We report that structure rather than four numbers that would look like tests.

V. Results

A. The nominal test is invalid; the calibrated one is adequately powered

Over 12,000 random-walk replicates spanning six lengths, the nominal one-sided t -test on θ_1 rejects 42.0% of the time rather than 5%. The inflation does not shrink with sample size. It grows, from 36.8% (95% CI [34.7, 38.9]) at $n = 100$ to 44.4% ([42.2, 46.5]) at $n = 1000$, which is the signature of spurious regression on an integrated series rather than a small-sample artefact. Surrogate calibration restores nominal size across the same grid: 5.7% pooled over 3,000 calibrated replicates, with every per-length interval covering 0.05. Calibration costs power, so we quote power for the test actually used. At $n = 200$ it is 1.00 for strong cusp dynamics, 0.76 for moderate and 0.24 for weak, reaching 1.00 in every regime by $n = 1000$ (Fig. 3). The null below therefore excludes strong and moderate dynamics. It does not exclude a faint fold.

B. The standard early-warning estimator fails on ground truth

Proposition 2 is a statement about the system. Testing it requires an estimator, and we checked ours before trusting it. Simulating from a model that genuinely contains a fold, the exponent computed from the true equilibrium curvature is +0.565 against a predicted +0.5, so theory and implementation agree. The rolling-window indicators that the early-warning literature relies on [15], [16] do not. Across a grid of window lengths, detrending choices

TABLE II

Can the standard early-warning estimator recover the exponent on ground-truth cusp data? The closed-form theory does. The rolling-window estimators return the wrong sign.

Estimator	Predicted	Detrend	Window	Median
theory $\lambda_{\text{rel}}(\mu)$	+0.5	—	—	+0.565
$-\log \text{AC1}$	+0.5	no	30	-0.262
$-\log \text{AC1}$	+0.5	no	120	-0.035
variance	-0.5	no	30	+0.311
variance	-0.5	yes	120	+0.036

and series lengths up to $n = 20,000$, rolling $-\log \text{AC1}$ returns negative slopes where +0.5 is predicted and rolling variance returns positive slopes where -0.5 is predicted (Table II).

The mechanism is structural, not statistical. Critical slowing down means the relaxation time diverges at the fold, so a fixed-width window saturates precisely where the effect is strongest, and detrending removes the low-frequency power carrying what remains. Raising n by an order of magnitude does not recover the sign. We therefore withdraw our own P2 test as uninterpretable, and the caution generalises: a directional early-warning result obtained this way is weak evidence either way.

C. The identifiable prediction is not supported

Under the calibrated test, the restoring cubic term is significant in 2.7% of the 147 units, against a nominal 5%. By corpus the figures are 0.0% for WESAD ($n = 15$), 0.0% for the examination sessions ($n = 30$) and 3.9% for the nurse sessions ($n = 102$). Because 102 units come from 15 nurses, we also weight one vote per person, which gives 1.5% (Wilson interval [0.0, 3.1%]), with 10% of people showing at least one significant session. That last figure is the most generous reading available and is still what chance produces across several sessions each.

Held-out prediction agrees from a different direction. The full model does not beat its own nested monostable special case: it wins on 33.8% of units, with a median difference of -0.25 nats per step and a Wilcoxon $p = 8 \times 10^{-5}$ in the simpler model's favour. Since the monostable model is this model with the cubic term switched off, that is the cleanest available statement that the cubic term is doing no work. Prospective onset prediction gives AUC between 0.50 and 0.46 at lead times of 3 to 12 windows, which is chance.

D. Structure without fold structure

The null is specific, and it would be wrong to summarise it as noise. Two of the three statistics we calibrate do exceed their surrogate nulls: the likelihood ratio against the monostable restriction at 40.8% of units, and the bimodality of the marginal distribution at 49.7%. Only the statistic that isolates the cubic term sits at chance. A bimodal marginal accompanied by no cubic drift is exactly the configuration in which a study reporting distributional

evidence alone would conclude the opposite of the correct answer.

Robustness is mixed, and we report it that way. Configuration B sits at nominal under both the random-walk (0.046) and the harder IAAFT (0.069) ensembles, and holds across 30s to 120s windows. Configurations A and C are elevated, at 0.11 to 0.16. Three explanations were tested and all failed: IAAFT did not bring C down, clipping rates run opposite to the effect, and a size study on AR(1) series (ϕ from 0 to 1) found no persistence at which the calibrated test over-rejects. What remains is nonlinear structure the surrogates do not reproduce. Since A is the negative control, built on a coordinate already shown to be noise-dominated, the parsimonious reading is artefact from motion and peak-detection failure rather than a fold. A test that fires on the negative control shows that significance does not mean cusp there, which is why the conclusion rests on B.

VI. Discussion and Conclusion

Modelling overload as motion in a cusp potential replaces a descriptive account with a generative one whose consequences are computable in advance. States become basins, the kernel becomes a function of nine parameters, and sudden onset, delayed recovery and two exponents become point hypotheses. The cost of that precision is that the model can lose, and here it did.

The identifiability result is the more durable contribution. Any attempt to fit cusp geometry to wearable physiology inherits the same ratio structure, so such a study needs either much longer records, on the order of $n \gtrsim 500$ windows, or a reparameterisation that estimates a and b directly. Reporting a splitting factor of 12 from a 200-window recording is not a finding, and the arithmetic that produces it is not specific to our implementation.

Limitations. No autistic participants are analysed, so the motivating claim is untested rather than refuted. Tonic electrodermal activity is a peripheral, artefact-prone proxy for a latent construct. Weak cusp dynamics remain compatible with our data. The estimator critique covers fixed-width rolling indicators with and without Gaussian detrending, which is what the applied literature mostly uses, and does not extend to adaptive-bandwidth or spectral estimators, which we did not test.

Toward validation. We pre-specify the study that would test the mechanism in its intended population: 28 days of ecological momentary assessment with continuous wrist physiology in autistic adults and a matched comparison group, using the same estimator, nulls and falsification criteria as here. The primary hypothesis is a between-group difference in α_0 , better powered than the within-group detection attempted here, and the design targets $n \gtrsim 500$ windows per participant so the control parameters sit inside the identifiable regime. A second well-powered null under those conditions would be evidence that overload is not a cusp catastrophe at physiological timescales.

Acknowledgment

This work uses the WESAD, Wearable Exam Stress and Nurse Stress datasets, and I thank their authors for releasing them openly. In accordance with IEEE policy on generative AI, the author discloses that an AI assistant was used for code implementation, literature organisation and manuscript drafting. The author verified all analyses, results and claims, and is solely responsible for the content.

References

- [1] K. MacLennan, S. O'Brien, and T. Tavassoli, "In our own words: The complex sensory experiences of autistic adults," *Journal of Autism and Developmental Disorders*, vol. 52, no. 7, pp. 3061–3075, 2022.
- [2] M. Kiep, A. Spek, E. Ceulemans, and I. Noens, "Sensory processing and executive functioning in autistic adults," *Journal of Autism and Developmental Disorders*, vol. 55, no. 6, pp. 2075–2084, 2025.
- [3] E. A. Demetriou et al., "Autism spectrum disorders: A meta-analysis of executive function," *Molecular Psychiatry*, vol. 23, no. 5, pp. 1198–1204, 2018.
- [4] N. Angeletos Chrysaitis and P. Seriès, "10 years of Bayesian theories of autism: A comprehensive review," *Neuroscience & Biobehavioral Reviews*, vol. 145, p. 105022, 2023.
- [5] M. Adamou, A. Kehagias, and G. Antoniou, "A multi-scale mathematical framework for modelling the dynamic nature of autism spectrum disorder symptoms: Integrating predictive coding, information theory, and network principles," *Frontiers in Psychiatry*, vol. 17, p. 1787120, 2026.
- [6] S. Cano, C. Cubillos, R. Alfaro, and A. Romo, "Wearable solutions using physiological signals for stress monitoring on individuals with autism spectrum disorder (ASD): A systematic literature review," *Sensors*, vol. 24, no. 24, p. 8137, 2024.
- [7] Y. Chen, Z. Xi, T. Greene, and W. Mandy, "A systematic review of ecological momentary assessment in autism research," *Autism*, vol. 29, no. 6, pp. 1374–1389, 2025.
- [8] P. D. Ditlevsen and S. J. Johnsen, "Tipping points: Early warning and wishful thinking," *Geophysical Research Letters*, vol. 37, no. 19, p. L19703, 2010.
- [9] C. Boettiger and A. Hastings, "Quantifying limits to detection of early warning for critical transitions," *Journal of the Royal Society Interface*, vol. 9, no. 75, pp. 2527–2539, 2012.
- [10] H. A. Kramers, "Brownian motion in a field of force and the diffusion model of chemical reactions," *Physica*, vol. 7, no. 4, pp. 284–304, 1940.
- [11] P. Schmidt, A. Reiss, R. Duerichen, C. Marberger, and K. Van Laerhoven, "Introducing WESAD, a multimodal dataset for wearable stress and affect detection," in *Proc. 20th ACM Int. Conf. Multimodal Interaction (ICMI)*, 2018, pp. 400–408.
- [12] M. R. Amin, D. Wickramasuriya, and R. T. Faghih, "A wearable exam stress dataset for predicting cognitive performance in real-world settings (version 1.0.0)," *PhysioNet*, 2022.
- [13] S. Hosseini et al., "A multimodal sensor dataset for continuous stress detection of nurses in a hospital," *Scientific Data*, vol. 9, p. 255, 2022.
- [14] T. Schreiber and A. Schmitz, "Improved surrogate data for nonlinearity tests," *Physical Review Letters*, vol. 77, no. 4, pp. 635–638, 1996.
- [15] M. Scheffer et al., "Early-warning signals for critical transitions," *Nature*, vol. 461, pp. 53–59, 2009.
- [16] V. Dakos et al., "Methods for detecting early warnings of critical transitions in time series illustrated using simulated ecological data," *PLOS ONE*, vol. 7, no. 7, p. e41010, 2012.